

1- get namespaces of the system

```
controlplane $ kubectl get namespaces
NAME                STATUS   AGE
default             Active   20h
kube-node-lease     Active   20h
kube-public         Active   20h
kube-system         Active   20h
local-path-storage  Active   20h
controlplane $
```

2- get amount of pods in the kube-system namespace

check

```
controlplane $ kubectl get pods -n kube-node-lease
No resources found in kube-node-lease namespace.
controlplane $
```

3- create a deployment of 2 replicas with the name of beta

```
Editor  Tab 1  +
apiVersion: apps/v1
kind: Deployment
metadata:
  name: beta
spec:
  replicas: 2
  selector:
    matchLabels:
      app: finance
  template:
    metadata:
      name: finance-redis
      labels:
        app: finance
    spec:
      containers:
        - name: beta
          image: redis
          resources:
            requests:
              cpu: 500m
              memory: 1G
            limits:
              cpu: 1
              memory: 2G
```

Creating the finance namespace

```

Editor  Tab1  +
apiVersion: v1
kind: Namespace
metadata:
  name: finance
~

```

Apply in the finance namespace

```
controlplane $ kubectl apply -f deploy.yml -n finance
```

```

controlplane $ vim finance.yml
controlplane $ kubectl get pods -n finance
NAME                                READY   STATUS    RESTARTS   AGE
beta-54d89cd7ff-m8pn9              1/1     Running   0           4m10s
beta-54d89cd7ff-sh66r              1/1     Running   0           4m10s
controlplane $

```

4- assigning label blue of master node

```

NAME                STATUS    ROLES    AGE   VERSION
controlplane        Ready     control-plane   21h   v1.31.0
node01              Ready     <none>         21h   v1.31.0
controlplane $ kubectl label nodes controlplane color=blue
node/controlplane labeled
controlplane $

```

5- creating deployment using node affinity on blue

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: blue
spec:
  replicas: 2
  selector:
    matchLabels:
      app: blue-app
  template:
    metadata:
      name: blue-app
      labels:
        app: blue-app
    spec:
      containers:
        - name: nginx-blue
          image: nginx
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            nodeSelectorTerms:
              - matchExpressions:
                  - key: color
                    operator: In
                    values:
                      - blue

```

6-creating iti namespace

```
apiVersion: v1
kind: Namespace
metadata:
  name: iti
```

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: iti-resource
  namespace: iti
spec:
  hard:
    pods: 2
```

```
resourcequota/iti-resource configured
controlplane $ kubectl get quota --namespace=iti
NAME           AGE    REQUEST    LIMIT
iti-resource   40s    pods: 0/2
```

7- creating nginx deployment with 3 replicas

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-app
  namespace: iti
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      name: finance-redis
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
```

8- how many pods were created and why

```
controlplane $ kubectl apply -f deployment.yml
deployment.apps/nginx-app created
controlplane $ kubectl get pods -n iti
NAME                                READY   STATUS             RESTARTS   AGE
nginx-app-696df79c4c-2nzjp         0/1     ContainerCreating   0           8s
nginx-app-696df79c4c-jpm5m         0/1     ContainerCreating   0           8s
controlplane $ kubectl get pods
No resources found in default namespace.
controlplane $ kubectl get deploy
No resources found in default namespace.
controlplane $ kubectl get deploy -n iti
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
nginx-app 2/3      2            2           28s
controlplane $ kubectl describe deploy nginx-app -n iti
Name:          nginx-app
Namespace:     iti
CreationTimestamp: Wed, 12 Feb 2025 15:57:16 +0000
Labels:
```

2 pods were created instead of the 3 specified by the deployment because the resource quota hard limit specified for the namespace are 2 pods only

9- get daemonsets on all namespaces

```
controlplane $ kubectl get daemonsets --all-namespaces
NAMESPACE   NAME      DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
kube-system  canal     2         2         2       2            2           kubernetes.io/os=linux 23h
kube-system  kube-proxy 2         2         2       2            2           kubernetes.io/os=linux 23h
controlplane $
```

10- what daemon sets exist on the kube-system

The canal daemon-set and the kube-proxy exists on all nodes

11- the image used by the kube-proxy daemonset

```
NAMESPACE   NAME      DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE   CONTAINERS
IMAGES
kube-system  canal     2         2         2       2            2           kubernetes.io/os=linux 23h  calico-node
annel        docker.io/calico/node:v3.24.1,quay.io/coreos/flannel:v0.15.1 k8s-app=canal
kube-system  kube-proxy 2         2         2       2            2           kubernetes.io/os=linux 23h  kube-proxy
registry.k8s.io/kube-proxy:v1.31.0
k8s-app=kube-proxy
```

registry.k8s.io/kube-proxy:v1.31.0

12- taint node 1

```
controlplane $ kubectl taint nodes node01 special-node=true:NoSchedule
node/node01 tainted
controlplane $
```

13- create a pod called tolerant pod

```

apiVersion: v1
kind: Pod
metadata:
  name: tolerant-pod
spec:
  containers:
    - name: nginx-pod
      image: nginx

```

14- which node is it deployed on?

```

controlplane $ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP          NODE          NOMINATED NODE   READINESS GATES
tolerant-pod   1/1     Running   0           31s   192.168.0.4 controlplane   <none>           <none>
controlplane $

```

Deployed on the controlplane node because the pod doesn't have a toleration that matches node1 so it's deployed on the controlplane node

15- deleting the pod and rescheduling

```

Editor: 1 tab
apiVersion: v1
kind: Pod
metadata:
  name: tolerant-pod
spec:
  tolerations:
    - key: "special-node"
      operator: "Equal"
      value: "true"
      effect: "NoSchedule"
  containers:
    - name: nginx-pod
      image: nginx

```

```

Error from server (NotFound): pods "tolerant-pod" not found
controlplane $ kubectl delete pod tolerant-pod
pod "tolerant-pod" deleted
controlplane $ kubectl apply -f tolerant-pod.yml
pod/tolerant-pod created
controlplane $ kubectl get pods -o wide
NAME          READY   STATUS             RESTARTS   AGE   IP          NODE    NOMINATED NODE   READINESS GATES
tolerant-pod   0/1     ContainerCreating   0           4s   <none>     node01   <none>           <none>
controlplane $

```

16- on which node is it scheduled

The pod is scheduled on the node01 because its tolerations now the node01 which are special-node=true:NoSchedule